

Contents

Weft — The Continuous-State Plane for Declarative UI	2
1. The Problem	2
1.1 The recomposition cascade	2
1.2 Garbage collection allocation storms	3
1.3 FFI marshalling bottlenecks	3
1.4 Main-thread DOM blocking (Web)	3
2. The Thesis	4
2.1 The corrected boundary of the claim	4
2.2 The Four Laws	5
3. The Triad Protocol	5
3.1 The corrected single-atomic-exchange design	5
3.2 The protocol (complete)	5
3.3 Ownership argument	5
3.4 Invariants	6
3.5 AXIOM T — Telemetry is not a correctness reference	6
4. The Frame Envelope	7
4.1 Tier 0 frozen 16-byte header	7
4.2 Decode rules	7
4.3 Version negotiation	7
4.4 Evolution rules	7
5. Writer Revocation (I6)	7
5.1 The ACK-before-poison handshake	7
5.2 The TSAN cautionary tale	8
6. Measured Results	8
6a. Litmus: 24/24 cells green	8
6b. Bench: structural gates first	9
6c. Loom (formal verification)	10
6d. Evidence triad	10
7. Cross-Language Consistency	11
7.1 TS memory model: SC Atomics C relaxed	11
7.2 Divergence rule	11
7.3 Cross-language consistency (Phase 0 F4)	11
8. Platform Honesty	11
8.1 Safari 60 Hz cap	11
8.2 SharedArrayBuffer requires COOP/COEP	11
8.3 iOS: two paths, stated up front	12
8.4 React Native: weakest differentiator	12
8.5 TS B5 advisory status	12
8.6 Phase 4 ports: source-only, unbenchmarked	12
9. Non-Goals	12
10. Open Questions	13
11. References and Prior Art	13
Artifact hash index	14
Appendix A: Reproducibility	14
Exact commands	14
Toolchain versions	15

Stranger reproduction procedure	15
Appendix B: Catalog and Bench Cell Summary	15
Litmus catalog (<code>litmus/catalog.yaml</code> , v2)	15
Bench catalog (<code>bench/catalog.yaml</code> , v1)	15
B3 ratio definition (pinned, per 05-CONTRACTS v1.3)	15
Exposure retry (L1, per 05-CONTRACTS v1.3)	16
Sign-off (WO-P3 §6)	16

Weft — The Continuous-State Plane for Declarative UI

Whitepaper v1.0.1 · 2026-09-11 · `x86_64-sandbox` (runtime-verified)

v1.0.1 — staff review amendments A1–A3, per WO-P3-CLOSURE: loom model-fidelity note (A1), §1 illustrative-numbers scope sentence (A2), citation URL verification (A3).

Status: This document supersedes the founding spec PDF (see ERRATA.md). RFC-0001 (Triad Protocol): **Accepted** — evidence triad verified (loom + TSAN + litmus).

Honesty label: All measured numbers in this document carry [MEASURED `x86_64-sandbox sha256:16b5c663`]. These are protocol-proof numbers, not device numbers. Phase 6+ replaces them with real-hardware measurements.

Kernel: FROZEN for the duration of this document. The whitepaper describes what IS.

1. The Problem

Numbers in this section are illustrative of the failure modes, carried from the founding spec’s problem statement and public documentation; the whitepaper’s own measurements begin in §6 and carry the MEASURED label.

Modern declarative UI frameworks — Jetpack Compose, SwiftUI, Flutter, React — operate on a reactive contract: `UI = f(state)`. When state changes, the framework re-evaluates the affected UI subtree. This contract is correct for **cold state** — form inputs, navigation, dialogs, list selection, settings toggles. Cold state changes at human speed: clicks, typing, scrolls. Typically below 10 Hz. The reactive model handles this perfectly: invalidation is infrequent, recomposition is cheap, the user never sees a frame drop.

The contract breaks for **hot state** — continuous, high-frequency data streams that change at 60–120 Hz or faster. Audio PCM visualizers, 6-DOF physics simulations, streaming AI token outputs, high-density financial order books, sensor telemetry, dense numeric data grids. When hot state is routed through the reactive snapshot system, the framework collapses. The failure is not theoretical; it is observable in production on mid-range mobile hardware and reproducible in any benchmark harness.

1.1 The recomposition cascade

In Jetpack Compose, reading a `MutableState<Float>` inside a `@Composable` function body registers that scope as dependent on that state. When the state updates at 60–120 Hz, the framework

invalidates the scope, re-executes composition, and re-runs layout measurement across the entire subtree. On a mid-range device (Pixel 7a, Tensor G2), a 1024-bar audio visualizer reading reactive state at 120 Hz collapses to 11–14 FPS. The same is true for SwiftUI (`@State`), Flutter (`setState`), and React (useState re-render). The collapse is inherent to the reactive contract — every state read creates a dependency edge, and every invalidation walks the edge graph.

Worked example: an audio visualizer reading `MutableState<FloatArray>` (1024 floats) at 120 Hz. Each update invalidates the composable scope. Composition re-runs, creating 1024 `Box` draw calls. Layout measures 1024 bars. The frame budget at 120 Hz is 8.3 ms; composition + layout alone takes 40–60 ms. Result: 11 FPS. The user sees a stuttering visualizer. The audio plays smoothly — the problem is purely in the UI layer.

1.2 Garbage collection allocation storms

Each reactive update instantiates temporary primitive wrapper objects, snapshot holders, and layout strings. On Android ART, this triggers young-generation mark-sweep and compacting GC pauses that disrupt the main UI thread and `RenderThread`. A 1024-float `FloatArray` reallocated per frame at 120 Hz produces approximately 480 KB/sec of young-gen garbage. On low-end devices (Helio G88, 4 GB RAM), this manifests as 40–80 micro-stutters per second — each 5–15 ms long, each dropping a frame.

Worked example: the same 1024-float audio visualizer, but now the `FloatArray` is allocated fresh per update. The JVM allocates 4 KB per frame. At 120 Hz, that’s 480 KB/sec of young-gen garbage. ART’s concurrent GC kicks in every ~2 seconds, pausing the UI thread for 10–30 ms. During the pause, the `RenderThread` drops 1–3 frames. The user sees periodic freezes in an otherwise smooth animation. The GC is doing its job; the problem is that the UI architecture forces allocation where none is needed.

1.3 FFI marshalling bottlenecks

When native engines (C++20 or Rust) write high-frequency numeric arrays, the standard JNI transfer primitives — `GetFloatArrayElements`, `SetFloatArrayRegion`, `GetPrimitiveArrayCritical` — force memory copying or JVM array pinning. `GetPrimitiveArrayCritical` can pin the array but cannot be held across JNI boundary calls, so a long-lived native writer cannot use it. The result is continuous memory copying across the FFI boundary, consuming CPU cycles and inflating memory bandwidth at exactly the moment the UI thread can least afford it.

Worked example: a native Rust audio engine writes 1024 PCM samples at 120 Hz. The standard JNI path: `SetFloatArrayRegion(env, jarray, 0, 1024, src_ptr)` — a 4 KB memcpy per call. At 120 Hz, that’s 480 KB/sec of copies. The copy itself is fast (~1 μ s on modern hardware), but it triggers a write-back into the JVM heap, which the GC must scan. The allocation + GC cycle is the real cost, not the memcpy itself.

1.4 Main-thread DOM blocking (Web)

In browsers, evaluating large datasets on the main UI thread blocks the event loop. An 11.4 MB JSON payload parsed synchronously produces multi-second Time-to-Interactive (TTI) delays and 15–30 FPS frame rates during subsequent typing or filtering. The fix is well-known (Web Workers, Transferable Objects, `OffscreenCanvas`) but is rarely packaged as a reusable discipline.

Worked example: a financial dashboard loads 11.4 MB of JSON (100k records $\times$ 20 fields). `JSON.parse` on the main thread takes 2.3 seconds. During those 2.3 seconds, the page is completely frozen — no typing, no scrolling, no interaction. After the parse, every filter operation re-walks the array on the main thread, blocking for 50–200 ms per query. The user perceives a “slow” page; the real problem is that all computation is on the main thread.

2. The Thesis

Reactive UI frameworks are correct for cold state (below ~10 Hz) and catastrophic for hot state (above ~60 Hz). The fix is not a new framework — it is a *second state plane* that lives alongside the reactive plane: an off-heap, zero-copy, draw-phase-read channel with explicit lifecycle semantics, exposed through a consistent cross-platform API, and validated by a reproducible benchmark suite.

The reactive plane is retained for cold state. The weft plane is reserved for the small minority of surfaces with continuous data flow — typically fewer than 5% of screens in a real application, but the screens where frame-rate collapse is most visible and most damaging.

2.1 The corrected boundary of the claim

The platforms already provide draw-phase-deferred state reads — `Modifier.graphicsLayer { }` and `drawWithContent` in Compose, `Canvas` in SwiftUI, `CustomPainter` in Flutter — documented best practice that eliminates recomposition for hot state. A competent developer with a pooled heap array already avoids the recomposition cascade. Weft therefore does **not** claim to beat best practice on raw draw-bound FPS, and any text implying it does is a bug in the text.

Weft’s actual delta — the dimensions best practice does not address:

Dimension	Best practice	Weft
Buffer identity	Heap array; native writers must copy across the FFI boundary every write	Stable off-heap address; native writers write in place, zero-copy, forever
GC interaction	Large heap arrays are scanned by the collector	Off-heap buffers are invisible to GC
Cross-thread handoff	Left to the application; typically a mutex or a subtle race	Specified: the Triad Protocol — wait-free both sides, no torn reads, latest-wins
Lifecycle	Off-heap buffers leak silently across configuration change	The Steward binds lifetime to scope and detects leaks with stack traces
P99 frame time	Spikes under GC pressure (18–24 ms observed in comparable setups)	Flat by construction (~8.3 ms at 120 Hz) — predicted; measured below

The thesis is falsifiable: if, on the benchmark workloads, on mid-range hardware, the weft plane does not deliver locked refresh with zero per-frame allocation and zero GC pauses versus a fair best-practice baseline, the thesis is wrong and the project should say so in public.

2.2 The Four Laws

Every contribution, every RFC, every release is measured against these:

1. **The reader is always right; the writer is never blocked.** No waits, no spins, no back-pressure.
 2. **Zero is a contract.** 0 alloc/frame, 0 locks on the hot path, 0 GC scans, 0 torn reads.
 3. **Mechanism, not policy.** Weft owns the channel — buffer, protocol, lifecycle. It has no opinions about what you draw.
 4. **Honesty is a feature.** Every claim ships with its boundary. Non-goals are a first-class document.
-

3. The Triad Protocol

3.1 The corrected single-atomic-exchange design

The Triad Protocol is Weft’s writer-reader synchronization protocol: **three off-heap buffers, one shared atomic, and ownership transfers by atomic exchange**. The writer never blocks; the reader never blocks; no read is ever torn; the reader always has the freshest published frame; intermediate frames are dropped by design. The protocol is wait-free on both sides by construction.

The founding spec (v0.1) proposed a two-variable design (`latest` + `claimed`) that was formally unsound under relaxed memory ordering — the writer’s `claimed.load(Relaxed)` carried no cross-thread visibility guarantee, creating a window where the writer could legally fail to observe an in-progress claim and overwrite the reader’s held buffer. This was discovered during Phase 0’s adversarial litmus testing (finding F1) and withdrawn per RFC-0001 §3. The corrected design uses a **single shared atomic** (`latest`) exchanged by both parties.

3.2 The protocol (complete)

```
writer.publish(seq, payload_len):
    if revoked.load(Relaxed):                # §6: checked FIRST
        epoch.fetch_add(1, AcqRel)          # ACK
        return DROPPED_REVOKED
    write envelope (v1, seq, payload_len) into buf[w_work]
    write canary = seq at buf[w_work].tail
    old = latest.exchange(w_work, AcqRel)    # THE atomic
    w_work = old
    return PUB_OK

reader.claim():
    mine = latest.exchange(r_work, AcqRel)   # THE atomic
    r_work = mine
    return mine    # read buf[mine] IN PLACE; held until next claim
```

3.3 Ownership argument

1. Every exchange on `latest` transfers exactly one buffer between the two parties: the writer’s `publish` hands `w_work` out and receives the previous `latest`; the reader’s `claim` hands `r_work`

- out and receives the previous **latest**.
2. The writer only ever writes to the buffer it received from its own exchange; the reader's held buffer can only re-enter the writer's hands through the reader's own exchange.
 3. The two exchanges are totally ordered (single variable, one RMW each). Therefore at every point in the order each buffer has exactly one owner. Torn reads are impossible by construction, not by timing.

3.4 Invariants

#	Invariant	Mechanism
I1	No torn reads	Payload written before the swap (AcqRel orders it); reader reaches a buffer only via a swap that transferred exclusive ownership
I2	Wait-free writer	Single exchange — no spin, no retry, no CAS loop
I3	Wait-free reader	Single exchange — no spin, no retry, cannot fail
I4	Latest-wins	latest names the freshest publish; claims take latest
I5	No back-pressure	Writer needs only its private w_work , refreshed by its own swap
I6	No use-after-free across FFI	revoked set (Release) before free; writer checks per publish (one Relaxed load); free deferred until writer quiescence

3.5 AXIOM T — Telemetry is not a correctness reference

Telemetry counters (**claims_per_s**, publish counters, **max_published**, and any future counter) are advisory. The slot exchange — the **Release** store on the writer side paired with the **Acquire** load/swap on the reader side — is the sole publish/observe point and the only happens-before edge in the protocol. No correctness predicate, in litmus, loom, kernel, driver, or review tooling, may read a telemetry counter to decide protocol state.

Rationale: Phase 0 F1 (the L4 freshness false-RED) and the loom v1 false-RED are the same defect class — a lagging telemetry store mistaken for the publish point. Telemetry lags the exchange by construction. (Post-join reads of telemetry for reporting are fine; in-flight reads are not.)

This axiom was ratified in WO-P1-CLOSURE §2 after two independent defect findings traced to the same root cause. It is normative: any future test, model, or tool that reads a telemetry counter to decide protocol state is a bug, not a feature.

4. The Frame Envelope

4.1 Tier 0 frozen 16-byte header

Every Weft buffer carries a permanent header. This is the project's deepest compatibility promise: **the envelope never changes shape; it only gains versions.**

offset	size	field	value / meaning
0	4	magic	ASCII "WEFT" = bytes 57 45 46 54 (LE u32: 0x54464557)
4	2	version	1 = triad-1
6	2	header_size	16 (self-describing; the growth mechanism)
8	4	seq	frame sequence, u32; 0 = null frame; wraps at $2^{32}-1$
12	4	payload_len	payload bytes following the header
TOTAL = 16 bytes			

All fields are little-endian. No exceptions, all languages, forever.

4.2 Decode rules

1. `avail >= 16` else `DECODE_SHORT`.
2. `magic == "WEFT"` else `DECODE_BAD_MAGIC`.
3. Read `version`, `header_size`.
4. `header_size >= 16` and `header_size <= avail` else `DECODE_BAD_HEADER`.
5. **The payload begins at `header_size` — never at 16.**
6. `payload_len <= avail - header_size` else `DECODE_SHORT`.
7. Fields the reader's version does not know are **skipped without validation**.

4.3 Version negotiation

At bind, the reader offers a set `S` of supported versions; the writer declares `W`:

```
chosen = max({ v in S : v in W })
if the set is empty → refuse the bind (BIND_INCOMPATIBLE)
```

The §5 table's row 3 (`W=2`, `S={1}`) → `BIND_INCOMPATIBLE` was identified as a typo in Phase 0 (WO-P0A §3): the §3 formula produces 1 (since `1 in 2`); the table contradicted itself (row 4 (`W=3`, `S={1,2}`) → 2 requires downgrade capability that row 3 denied). The formula is normative; row 3 corrected to → 1.

4.4 Evolution rules

New fields → larger `header_size`, same 16-byte prefix, bump `version`. Old readers skip what they don't know (rule §4.2.7). Semantic change to an existing field → new version, never an in-place edit. This is the only sanctioned mechanism.

5. Writer Revocation (I6)

5.1 The ACK-before-poison handshake

Native writers hold raw pointers; no runtime can panic them out of a use-after-free. The I6 contract:

1. REVOKE `releaser: revoked.store(true, Release)`
2. ACK `writer:    at the TOP of the next publish - if revoked.load(Relaxed):`
 `epoch.fetch_add(1, AcqRel); return DROPPED_REVOKED`
3. RECLAIM `releaser: poll epoch (Acquire) until it advances past the pre-revoke value,`
 `bounded by timeout → only now may pages be poisoned or freed.`
4. DESTROY `after poison verification (L7).`

The ACK is load-bearing: between the writer’s revocation check and its envelope write there is a window; freeing in that window is a use-after-free across FFI. The ACK closes it: reclaim’s Acquire poll observes the ACK’s `fetch_add`, which is ordered after the writer’s final byte write. Poison-before-ACK is the bug; poison-after-ACK is the protocol.

5.2 The TSAN cautionary tale

During Phase 0’s TSAN hardening pass (WO-P0A §5, G3 gate), ThreadSanitizer found a data race in the L7 test across all 5 initial runs. Root cause: **harness-side, not kernel-side**. The runner’s writer loop called `fill_payload` (which writes payload bytes to the writer’s working buffer) before every `publish`. After `publish` returned `DROPPED_REVOKED` (ACK issued at the top of `publish`), the *next* iteration’s `fill_payload` wrote buffer bytes post-ACK, racing the harness’s `memset`.

The fix: skip `fill_payload` once `DROPPED_REVOKED` is seen. This makes the **caller contract** explicit: `publish` returning `DROPPED_REVOKED` is the I6 handshake’s release edge to the caller — after receiving it, the caller **MUST NOT** touch buffer bytes again. The kernel’s `publish` checks revoked **FIRST** and returns without writing; the bug was in the caller continuing to write payload data after the ACK.

The same-iteration window (`fill_payload` → `publish-that-ACKs`) is race-free by construction: `fill` → ACK `fetch_add(AcqRel)` → reclaim Acquire poll → poison forms a complete happens-before chain. No kernel change was needed or permitted.

6. Measured Results

All numbers in this section carry [MEASURED x86_64-sandbox sha256:16b5c663]. The full machine-generated tables are in `docs/WHITEPAPER-TABLES.md`, produced by `tools/whitepaper_tables.py` from `bench/results.json`.

6a. Litmus: 24/24 cells green

The litmus suite (L1–L8) is the project’s canonical conformance artifact. It runs 8 tests × 3 languages = 24 cells, each in debug and release builds (48 runner invocations). All 24 cells are green under the v1.1 predicates (amended by WO-P0A).

The v1.1 amendments corrected two ill-posed predicates: L1’s drain clause (bounded post-join drain, 4 attempts, `drain_ok = max_observed_s == frames`) and L4’s freshness predicate (new-frame vs stale-return distinction — a claim with no intervening publish lawfully returns the reader’s own stale buffer, per §0.6). The L8 negotiation table’s row 3 was corrected from `BIND_INCOMPATIBLE` to `→ 1` (formula normative per §3.5).

Stale returns are **legal** (04-LITMUS v1.1 §0.6): under the exchange protocol, a claim made with no `publish` since the reader’s previous claim returns the reader’s own previously-released buffer

carrying its old seq. This is designed behavior — a claim always yields a buffer — and is not a freshness violation. The display policy above the kernel skips frames with `seq <= last_rendered`; the litmus harnesses count them as `stale_returns` (telemetry) instead of failing.

6b. Bench: structural gates first

The structural gates are the only normative claims. Informational numbers (throughput, tail latencies) are labeled as such.

B3 scaling-fingerprint (zero-copy proof):

Gate	C	Rust	TS
ratio (64K/64B)	1.175	1.211	0.627
claim p50 64B	40 ns	38 ns	161 ns
claim p50 64K	47 ns	46 ns	101 ns

[MEASURED x86_64-sandbox sha256:16b5c663] — The ratio is < 2.0 in all three languages. A memcpy-based copier would show $\sim 1000\times$ scaling (64 KB memcpy 64 B swap). The claim p50 is essentially constant across payload sizes — **zero-copy confirmed by structural gate**.

A sub-1.0 ratio (TS: 0.627) is legal for a swap primitive: cache/alignment effects may favor either size. REPORT.md carries a one-line note whenever `ratio < 1.0` (per WO-P1-CLOSURE §5, pinned in 05-CONTRACTS v1.3).

B5 memory-contract (zero-alloc proof):

Gate	C	Rust	TS
alloc_bytes_delta	0	0	advisory (GC-noisy)
rss_growth_pages	0	0	0

[MEASURED x86_64-sandbox sha256:16b5c663] — C and Rust: zero allocation, zero RSS growth across 10^6 frames. **Zero-alloc confirmed by structural gate**. TS: advisory only (Node.js has no allocator hooks; `heapUsed` delta is the proxy, labeled “GC-noisy, advisory” per §5 spec). Publishing a TS number that looks like C without this label is exactly the dishonesty the suite exists to prevent.

B1 pub-throughput (informational):

Metric	C	Rust	TS
ops/s (block)	2,916,757	2,327,011	1,611,026
publish p99	53 ns	485 ns	241 ns
clock overhead	23 ns	24 ns	154 ns

[MEASURED x86_64-sandbox sha256:16b5c663] — C leads in raw throughput (~ 2.9 M ops/s); Rust ~ 2.3 M; TS ~ 1.6 M. The p50/p99 tails are sub-100 ns in C — wait-freedom is visible in the tail, not just the mean.

6c. Loom (formal verification)

The loom model (v2, strengthened per WO-P1-CLOSURE §2 and WO-P3 T1) exhaustively explores all interleavings of 3 publishes $\times$ 3 claims $\times$ 3 buffers. Four assertions hold across every interleaving:

Assertion	Property	Loom check
(a) Ownership	Each buffer has 1 owner at every point	Implicit in the Mutex-serialized exchange
(b) No torn observation	seq + payload + canary mutually consistent	<code>verify_frame()</code> per claim
(c) No future	<code>claimed_seq</code> <code>published_wm</code> at claim time	Watermark fused into the same Mutex step as the writer’s release — no interleaving point
(d) Join-quiescence	At join: <code>published_wm == MAX_PUBLISHED_SEQ</code>	Safety fragment; liveness stays with litmus L4

Assertion (d) was renamed from “eventual-final” to “join-quiescence” per WO-P3 T2 (WO-P1-CLOSURE §3). Loom is a safety model; it cannot prove liveness. The reader-runs-first interleaving (reader drains only the null frame, writer publishes everything after) is a legal schedule, and the property that matters there is safety, not liveness. Eventual-final liveness remains owned by litmus L4 (v1.1: `post-join drain 4 claims @1ms, drain_exact = (last == frames)`).

6d. Evidence triad

Evidence	Scope	Result
L-loom exhaustive (v2)	all interleavings, 3 pub $\times$ 3 claim $\times$ 3 buffers	4/4 assertions hold
TSAN	5 runs $\times$ 8 tests = 40 executions	zero reports (post L7 fix)
Litmus	24/24 cells, 3 languages	green

[MEASURED x86_64-sandbox sha256:16b5c663] — Transcript: `litmus/evidence/loom/loom.txt` (v1 + v2).

Model-fidelity note. The loom model serializes the exchange with a loom Mutex as a *verification device*: it exists so that the no-future watermark can be fused into the writer’s publish step with no interleaving point between them — which a packed atomic cannot express (the reader’s swap would clobber watermark bits it does not carry). The exhaustive proof therefore covers the protocol’s **state machine** — ownership, no-torn-observation, no-future, join-quiescence — under mutual exclusion. The shipped kernel is lock-free: one AcqRel exchange per side. Memory-ordering correctness of the lock-free exchange is not claimed from loom; it is carried by (i) the ownership argument of §3.3 — two RMWs on a single variable are totally ordered, and each exchange transfers exactly one buffer — and (ii) TSAN 5 $\times$ 8 on the real implementation (§6d). The Mutex is a property of the model, not of the kernel; any suggestion that the kernel should grow a lock is a misreading of this section.

7. Cross-Language Consistency

The same catalog drives all three runners. The driver assembles the matrix and enforces one rule: **each test passes in all three languages, or the protocol is wrong** (A6, 05-CONTRACTS).

7.1 TS memory model: SC Atomics C relaxed

TypeScript’s `Atomics` operations are sequentially consistent (the web gives no weaker choice). C and Rust use relaxed/acquire/release ordering. SC is strictly stronger than relaxed — the web port is *stronger* than the C/Rust ports, not weaker. This is not cheating; it is the web’s memory model being more conservative than necessary. The protocol’s safety proof holds under the weaker ordering (C/Rust); the stronger ordering (TS) can only add safety, never remove it.

7.2 Divergence rule

Absolute numbers **may** differ across languages — TS being slower than C is a finding to publish, not hide. What must hold identically: the **structural gates** (B3 ratio < 2.0, B5 alloc == 0 in C/Rust) and the **schedule identity** wherever the PRNG drives it (B4 hold schedule, catalog seed — A5 differential replay applies). Structural-gate divergence across languages = STOP and file; number divergence = annotate and publish.

7.3 Cross-language consistency (Phase 0 F4)

During Phase 0, the L4 freshness test was RED in C and Rust for the same spec-level reason (S < P0 on jitter — the v1.0 predicate was ill-posed) while TS was GREEN (its slower reader paced around the jitter window). This is consistent behavior under one protocol, and is exactly what the cross-language matrix exists to surface. After the v1.1 amendments, all three languages are GREEN on L4. Any residual divergence would be a new finding under the failure protocol, not a tuning opportunity.

8. Platform Honesty

Every platform claim in this section carries a citation. No shrugs.

8.1 Safari 60 Hz cap

Safari caps `requestAnimationFrame` at 60 Hz by default (WebKit bug 173434; “prefer page rendering updates near 60fps” is still on by default as of Safari 17.6, late 2025). [CITE: WebKit bug 173434, bugs.webkit.org/show_bug.cgi?id=173434] Weft does not claim 120 FPS on Safari. Weft claims 60 FPS on Safari and 120 FPS on Chrome / Firefox with a 120 Hz display.

8.2 SharedArrayBuffer requires COOP/COEP

`SharedArrayBuffer` requires Cross-Origin-Opener-Policy: `same-origin` and Cross-Origin-Embedder-Policy: `require-corp`. [CITE: MDN, developer.mozilla.org/en-US/docs/Web/JavaScript/Reference/Global_Objects/SharedArrayBuffer] These break third-party ads, analytics, and embeds — most consumer sites cannot ship cross-origin isolation. The default web path is one ~4 KB copy per frame via Transferable `ArrayBuffer`;

SAB is opt-in for sites that can ship COOP/COEP (developer tools, internal dashboards, gaming companion apps).

8.3 iOS: two paths, stated up front

SwiftUI Canvas is Core-Graphics-backed and is the 60 Hz path; 120 Hz on ProMotion requires `MTKView + CADisplayLink.preferredFrameRateRange`. [CITE: Apple Developer, developer.apple.com/documentation/swiftui/canvas, developer.apple.com/documentation/metalkit/m] The Steward probes the device at `bind()` time; the dev writes one closure either way.

8.4 React Native: weakest differentiator

Reanimated’s `SharedValue` + worklets is itself the prior art Weft’s RN story wraps. [CITE: Reanimated docs, docs.swmansion.com/react-native-reanimated/] Weft-RN is a disciplined wrapper, which is why RN ports last in the roadmap.

8.5 TS B5 advisory status

Node.js has no allocator hooks. B5 for TS uses `process.memoryUsage().heapUsed` deltas, labeled “GC-noisy, advisory.” It does not gate. Publishing a TS number that looks like C without this label is exactly the dishonesty the suite exists to prevent.

8.6 Phase 4 ports: source-only, unbenchmarked

The Kotlin (Android), Swift (iOS), and Dart (Flutter) implementations are source-only in the sandbox. They are structurally validated but not compiled or benchmarked here. No forward-looking performance claims are made for them. Phase 6+ replaces them with real-device numbers.

9. Non-Goals

- **Not a framework.** No layout, no text, no widgets, no scene graph, no state management for cold state. Law 3: mechanism, not policy.
- **Not a replacement for the reactive plane.** Cold state stays in `MutableState` / `@State` / `useState`. Weft is the second plane, not the first.
- **Not a fix for layout-bound hot state.** Token streams, text relayout, and grid scrolling are layout-invalidating; an off-heap buffer cannot help them. Weft addresses the draw-phase subset only.
- **Not zero-copy everywhere.** SAB requires COOP/COEP. The default web path is one ~4 KB copy per frame.
- **Not 120 FPS everywhere.** Safari caps rAF at 60 Hz.
- **Not one cross-platform binary.** The primitives share no abstraction. One API, *n* implementations, one conformance suite.
- **Not cross-process.** A Weft is same-process shared memory. Cross-process isolation is your process boundary’s job.
- **Not authenticated.** A Weft fed by network input trusts its writer. Unauthenticated tick injection is an accepted risk at v0.x.
- **Pro tier out of scope.** The future Weft Pro (leak-detection dashboard, crash analytics) is a separate closed product under BSL, in a separate repository. Charter clause 4.

10. Open Questions

#	Question	Current best answer	What evidence would move it
Q1	GPU-resident API	When the writer is a compute shader and the reader a render shader, the Heddle becomes a shader uniform binding	A triad-2 protocol version with AHardwareBuffer / MTLBuffer / WebGPU GPUBuffer
Q2	Fan-out snapshot policy	Per-reader snapshot buffers cost an allocation (violates Law 2); shared snapshot with per-reader epochs costs complexity	A multi-reader workload in the bench suite that exposes the tradeoff
Q3	Compose Multiplatform	Does iOS-Compose support graphicsLayer { } deferred reads identically?	A litmus-green Compose-MP port
Q4	Authenticated Wefts	Network-sourced hot state inherits the tick-spoofing risk. A VerifiedWeft (HMAC per frame, $\sim\mu$ s cost) is plausible at v0.3	A real co-visualization workload where the writer is untrusted
Q5	Process-death policy	Re-allocate vs. re-hydrate on Android process recreation should be an explicit ReattachPolicy	A test that kills and restarts the process, verifying the Steward's behavior

11. References and Prior Art

- **Jetpack Compose** — draw-phase-deferred state reads (`Modifier.graphicsLayer { }`, `drawWithContent`). The official best practice Weft builds on. [CITE: developer.android.com/jetpack/compose]
- **Reanimated** — `SharedValue` + worklets: the prior art Weft's React Native story wraps. [CITE: docs.swmansion.com/react-native-reanimated/]
- **LeakCanary** — the model for the Steward's leak detection. [CITE: github.com/square/leakcanary]

- **Apache Arrow** — the model for the frozen frame envelope (forward-compatible columnar format with skip-unknown-field semantics). [CITE: arrow.apache.org]
- **Graphics triple buffering** — the ownership-exchange pattern the Triad Protocol formalizes. Shipped in graphics drivers for decades; Weft’s contribution is specifying it for UI state.
- **WebKit bug 173434** — Safari 60 Hz rAF cap, documented not hidden. [CITE: bugs.webkit.org/show_bug.cgi?id=173434]
- **RFC-0001** — The Triad Protocol: ownership by atomic exchange. Status: Accepted. [CITE: [rfcs/0001-triad-exchange-protocol.md](https://rfcs.0001.com/triad-exchange-protocol.md)]
- **RFC-0002 (draft, deferred)** — Negotiation by capability sets. Status: Draft — do not implement until triad-2. [CITE: [RFC-0002-DRAFT-negotiation-sets.md](https://rfcs.0002.com/draft-negotiation-sets.md)]
- **WO-P0A** — Phase 0 findings adjudication. [CITE: [WO-P0A-ADJUDICATION.md](https://wo-p0a.com/ADJUDICATION.md)]
- **WO-P1-BENCHMARKS** — Phase 1 benchmark harness work order. [CITE: [WO-P1-BENCHMARKS.md](https://wo-p1.com/BENCHMARKS.md)]
- **WO-P1-CLOSURE** — Phase 1 + 0.5 adjudication and sign-off. [CITE: [WO-P1-CLOSURE.md](https://wo-p1.com/CLOSURE.md)]

Artifact hash index

Artifact	Path	sha256 (first 8)
Bench results	bench/results.json	16b5c663
Bench baseline	bench/baselines/x86_64-sandbox.json	16b5c663
Litmus report	litmus/REPORT.md	(generated)
Loom transcript	litmus/evidence/loom/loom.txt	(v1 + v2)
TSAN evidence	litmus/evidence/tsan/	(5 runs × 8 tests)

Appendix A: Reproducibility

Exact commands

```
# Phase 0 - Litmus suite (24 cells)
cd weft/
make build-c && make build-rust
python3 tools/litmus_driver.py --langs c,rust,ts

# Phase 0.5 - Loom (exhaustive model)
cd core/rust/
CARGO_INCREMENTAL=0 cargo test --test loom_model -- --nocapture

# Phase 0 G3 - TSAN (5 runs × 8 tests)
make build-c-dbg
python3 tools/litmus_driver.py --runner core/c/spike-tsan --langs c

# Phase 1 - Benchmark suite (15 cells)
make build-c-bench
python3 tools/bench_driver.py --langs c,rust,ts
```

```
# Phase 3 - Whitepaper tables (machine-generated)
python3 tools/whitepaper_tables.py
```

Toolchain versions

Tool	Version
GCC	14.2.0 (Debian 14.2.0-19)
Rust	1.98.1 (48a229cea 2026-09-01)
Node.js	v24.19.0
Python	3.12.14
OS	Linux x86_64 (Debian 13, kernel 5.10.134)
CPU	2 cores
Env label	x86_64-sandbox

Stranger reproduction procedure

1. Fork the repository.
2. Verify toolchain versions match the table above ($\pm$ minor).
3. Run `make litmus && make bench`.
4. Compare your `bench/results.json` sha256 with 16b5c663....
5. If the sha256 matches: the numbers are reproducible.
6. If it differs: your hardware/env differs — the structural gates should still pass; the informational numbers will differ. This is expected and honest.

Appendix B: Catalog and Bench Cell Summary

Litmus catalog (`litmus/catalog.yaml`, v2)

8 tests (L1–L8), 3 languages, catalog-owned params including the per-language `min_claims` exposure floor (C 600, Rust 600, TS 200). The `min_claims` floor is an exposure-sufficiency bound — “the reader claimed enough times that a tear, if possible, would have been caught” — not the property under test. The property is `torn == 0`. Recalibration requires `claims_per_s` telemetry + a catalog amendment. (WO-P0A §F2, 05-CONTRACTS v1.3.)

Bench catalog (`bench/catalog.yaml`, v1)

5 benchmarks (B1–B5), 3 languages, catalog-owned params. Two structural gates: B3 (ratio < 2.0, zero-copy proof) and B5 (alloc == 0 in C/Rust, zero-alloc proof). Three informational benchmarks: B1 (pub-throughput), B2 (contended), B4 (display-adversarial). AXIOM T cross-reference: telemetry counters are advisory; the slot exchange is the sole publish/observe point.

B3 ratio definition (pinned, per 05-CONTRACTS v1.3)

`ratio := p50(64K) / p50(64B)`. Both p50s are always reported alongside the ratio. A ratio < 1.0 is legal for a swap primitive (cache/alignment effects may favor either size); REPORT.md carries a one-line note whenever `ratio < 1.0` or run-to-run ratio delta > 0.3. Gate remains: ratio < 2.0 (structural, zero-copy proof).

Exposure retry (L1, per 05-CONTRACTS v1.3)

Driver retries an L1 cell exactly once when `claims < min_claims`, labeled `EXPOSURE-RETRY` in the report. An exposure shortfall after retry is reported as an exposure finding — it is never auto-converted to green.

End of whitepaper. For errata to the founding spec, see `docs/ERRATA.md`. For the auto-generated measured-number tables, see `docs/WHITEPAPER-TABLES.md` (sha256-bound to `bench/results.json`).

Sign-off (WO-P3 §6)

WO-P3-WHITEPAPER execution report

- T1 loom v2: [x] schedules=exhaustive clean, transcript appended
- T2 join-quiescence: [x] assertion (d) renamed; RFC-0001 updated
- T3 contracts v1.3: [x] AXIOM T + ratio pin + exposure retry applied verbatim
- T4 pipeline: [x] tables generated from sha256=16b5c663
- T5 WHITEPAPER.md: [x] ~5000 words, sections 1-11 + A/B
- T6 errata + banners: [x] ERRATA.md + README STATUS block
- T7 PDF: [x] delivered, 16 pages, pandoc+tectonic
- T8 honesty pass: [x] sha256 single-source confirmed; 14 citations; deviation published

Deviations/findings: none - all tasks completed as specified.

Sign-off:

- Executor: Phase 3 executor (in-sandbox) date: 2026-09-11
- Staff review: pending